

Analyze A/B Test Results

Table of Contents

- [Introduction](#)
- [Part I - Probability](#)
- [Part II - A/B Test](#)
- [Part III - Regression](#)

Part I - Probability

To get started, let's import our libraries.

```
import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. **Use your dataframe to answer the questions in Quiz 1 of the classroom.**

a. Read in the dataset and take a look at the top few rows here:

```
#first we will import the data from csv file
df=pd.read_csv("ab_data.csv")
#showing first 10 rows
df.head(10)
```

	user_id	timestamp	group	landing_page
converted				
0	851104	2017-01-21 22:11:48.556739	control	old_page
0				
1	804228	2017-01-12 08:01:45.159739	control	old_page
0				
2	661590	2017-01-11 16:55:06.154213	treatment	new_page
0				
3	853541	2017-01-08 18:28:03.143765	treatment	new_page
0				
4	864975	2017-01-21 01:52:26.210827	control	old_page
1				
5	936923	2017-01-10 15:20:49.083499	control	old_page
0				
6	679687	2017-01-19 03:26:46.940749	treatment	new_page
1				

```

7    719014    2017-01-17 01:48:29.539573    control    old_page
0
8    817355    2017-01-04 17:58:08.979471    treatment    new_page
1
9    839785    2017-01-15 18:11:06.610965    treatment    new_page
1

```

b. Use the below cell to find the number of rows in the dataset.

```

#finding the number of rows in data
df.shape[0]

```

294478

c. The number of unique users in the dataset.

```

#find how many unique users
df.user_id.nunique()

```

290584

d. The proportion of users converted.

```

#the porportion of users converted
df['converted'].mean()

```

0.11965919355605512

e. The number of times the new_page and treatment don't line up.

```

#we will metion new_page as a
#and mention treatment as b
a_with_no_b=df.query("landing_page == 'new_page' & group !=
'treatment'")
b_with_no_a=df.query("landing_page != 'new_page' & group ==
'treatment'")
#print(len(a_with_no_b))
#print(len(b_with_no_a))
len(a_with_no_b) + len(b_with_no_a)

```

3893

f. Do any of the rows have missing values?

```

#check if ew have missing values
df.isnull().sum().any()

```

False

2. For the rows where **treatment** is not aligned with **new_page** or **control** is not aligned with **old_page**, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to provide how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

```
df1 = df.drop(df[(df.group == "treatment") & (df.landing_page != "new_page")].index)
```

```
df2 = df1.drop(df1[(df1.group == "control") & (df1.landing_page != "old_page")].index)
```

```
<ipython-input-256-fd0a62ac880b>:3: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
```

```
df2 = df1.drop(df1[(df1.group == "control") & (df1.landing_page != "old_page")].index)
```

```
df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].shape[0]
```

```
0
```

3. Use **df2** and the cells below to answer questions for **Quiz3** in the classroom.

a. How many unique **user_ids** are in **df2**?

```
df2.user_id.nunique()
```

```
290584
```

b. There is one **user_id** repeated in **df2**. What is it?

```
#There is one user_id repeated in df2
```

```
df2.user_id.duplicated().sum()
```

```
#What is it
```

```
df2[df2['user_id'].duplicated()==True]['user_id']
```

```
2893    773192
```

```
Name: user_id, dtype: int64
```

c. What is the row information for the repeat **user_id**?

```
#entire row information about duplicated row
```

```
df2[df2.duplicated('user_id',keep =False)]
```

	user_id	timestamp	group	landing_page
converted				
1899	773192	2017-01-09 05:37:58.781806	treatment	new_page
0				
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page
0				

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

```
#remove one of tow duplicated rows
```

```
df2.drop(labels=2893,inplace=True)
```

```
#checking if the row removed or no  
df2[df2.duplicated('user_id', keep =False)].shape[0]
```

0

4. Use **df2** in the below cells to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

```
#poportion of converted in new data  
df2['converted'].mean()
```

0.11959708724499628

b. Given that an individual was in the control group, what is the probability they converted?

```
#porpotinon of converted which its group is control  
df[df['group']=='control']['converted'].mean()
```

0.12039917935897611

c. Given that an individual was in the treatment group, what is the probability they converted?

```
#porpotinon of converted which its group is treatment  
df[df['group']=='treatment']['converted'].mean()
```

0.11891957956489856

d. What is the probability that an individual received the new page?

```
#porpotinon of landing page is new page  
len(df[df['landing_page']=='new_page'])/len(df['landing_page'])
```

0.5

e. Consider your results from a. through d. above, and explain below whether you think there is sufficient evidence to say that the new treatment page leads to more conversions.

No, there is no sufficient evidence

The test seems to be well designed. Half of the population received the old_page and half of the population received the new_page.

12.04% that received the old_page were converted. 11.89% that received the new_page were converted. In conclusion, the new_page did not increase the conversion rate.

Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of p_{old} and p_{new} , which are the converted rates for the old and new pages.

Put your answer here.

2. Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

Use a sample size for each page equal to the ones in **ab_data.csv**.

Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the **convert rate** for p_{new} under the null?

```
#we will use the entire dataset to compute the conversion
#rate of new page because the null hypotheses consider that there is
#no difference between old and new page conversion
```

```
p_new=df2['converted'].mean()
p_new
```

0.11959708724499628

b. What is the **convert rate** for p_{old} under the null?

```
#we will use the entire dataset to compute the conversion
#rate of old page because the null hypotheses consider that there is
#no difference between old and new page conversion
```

```
p_old=df2['converted'].mean()  
p_old
```

0.11959708724499628

c. What is n_{new} ?

```
#number of users who landing new page  
n_new=len(df2[df2['landing_page']=='new_page'])  
n_new
```

145310

d. What is n_{old} ?

```
#number of users who landing old page  
n_old=len(df2[df2['landing_page']=='old_page'])  
n_old
```

145274

e. Simulate n_{new} transactions with a convert rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

```
new_page_converted = np.random.binomial(1,p_new,n_new)  
new_page_converted.mean()
```

0.12085885348565137

f. Simulate n_{old} transactions with a convert rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

```
old_page_converted = np.random.binomial(1,p_old,n_old)  
old_page_converted.mean()
```

0.1197048336247367

g. Find $p_{new} - p_{old}$ for your simulated values from part (e) and (f).

```
new_page_converted.mean() - old_page_converted.mean()
```

-0.00010535664448757531


h. Simulate 10,000 $p_{new} - p_{old}$ values using this same process similarly to the one you calculated in parts **a. through g.** above. Store all 10,000 values in a numpy array called **p_diffs**.

```
# Simulate 10,000 pnew - pold  
p_diffs = []  
for _ in range( 10000):  
  
    new_page_converted = np.random.binomial(1,p_new,n_new).mean()
```

```
old_page_converted = np.random.binomial(1,p_old,n_old).mean()
p_diffs.append(new_page_converted - old_page_converted)
```

i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

```
#showing the histogram of the p_diffs.
plt.hist(p_diffs);
plt.xlabel('p_diffs')
plt.ylabel('the Frequency')
plt.title(' Difference of new_page & old_page ');
```



j. What proportion of the **p_diffs** are greater than the actual difference observed in **ab_data.csv**?

```
df_control = df2[df2['group'] == "control"]
df_treatment = df2[df2['group'] == "treatment"]
# display the observed difference in ab_data
observed_diffs = df_treatment.converted.mean() -
df_control.converted.mean()
observed_diffs
```

```
-0.0015782389853555567
```

```
#proportion of difference between p_diffs and observed_diffs
(p_diffs > observed_diffs).mean()
```

91% is proportion of the p_diffs are greater than the actual difference observed in ab_data

k. In words, explain what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

this value is called p value and it means that we cannot reject the null hypothesis and we don't have an evidence that the new_page has a higher conversion rate than the old_page.

l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let n_old and n_new refer to the number of rows associated with the old page and new pages, respectively.

```
import statsmodels.api as sm

n_new = len(df2[df2['landing_page'] == "new_page"])
n_old = len(df2[df2['landing_page'] == "old_page"])
convert_new = len(df2.query('landing_page == "new_page" & converted == 1'))
convert_old = len(df2.query('landing_page == "old_page" & converted == 1'))
print(' convert_new is :', convert_new, '\n',
      'convert_old is :', convert_old, '\n',
      'n_new is :', n_new, '\n',
      'n_old is :', n_old)

convert_new is : 17264
convert_old is : 17489
n_new is : 145310
n_old is : 145274
```

m. Now use stats.proportions_ztest to compute your test statistic and p-value. [Here](#) is a helpful link on using the built in.

```
#computing our statistics with built_in function (statsmodels)
z_score, p_value = sm.stats.proportions_ztest([convert_old,
convert_new], [n_old, n_new], value=None, alternative='smaller',
prop_var=False)

print(' z score is : ', z_score, '\n', 'p value is : ', p_value)

z score is : 1.3109241984234394
p value is : 0.9050583127590245
```


n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts **j.** and **k.**?

z-score and p_value mean that one doesn't reject the Null.

The Null hypothesis mean the converted rate of the old_page is equal to or greater than the converted rate of the new_page.

The p_value is 0.91 and is higher than 0.05 significance level which is premitted. That means we can not be confident with a 95% confidence level that the converted rate of the new_page is larger than the old_page.

yes its agree (p value is 91% in all result)

Part III - A regression approach

1. In this final part, you will see that the result you acheived in the previous A/B test can also be acheived by performing regression.

a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

it should be logistic regression because the dependant variable is a binary variable

b. The goal is to use **statsmodels** to fit the regression model you specified in part **a.** to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create a column for the intercept, and create a dummy variable column for which page each user received. Add an **intercept** column, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

```
df2[['c', 'treatment']] = pd.get_dummies(df2['group'])
```

```
df2 = df2.drop('c', axis = 1)
df2.head()
```

	user_id	timestamp	group	landing_page
0	851104	2017-01-21 22:11:48.556739	control	old_page
1	804228	2017-01-12 08:01:45.159739	control	old_page
2	661590	2017-01-11 16:55:06.154213	treatment	new_page
3	853541	2017-01-08 18:28:03.143765	treatment	new_page
4	864975	2017-01-21 01:52:26.210827	control	old_page

	treatment
0	0
1	0
2	1
3	1
4	0

```
df2.rename(columns={'treatment': 'ab_page'}, inplace=True)
df2
```

	user_id	timestamp	group	landing_page \
0	851104	2017-01-21 22:11:48.556739	control	old_page
1	804228	2017-01-12 08:01:45.159739	control	old_page
2	661590	2017-01-11 16:55:06.154213	treatment	new_page
3	853541	2017-01-08 18:28:03.143765	treatment	new_page
4	864975	2017-01-21 01:52:26.210827	control	old_page
...
294473	751197	2017-01-03 22:28:38.630509	control	old_page
294474	945152	2017-01-12 00:51:57.078372	control	old_page
294475	734608	2017-01-22 11:45:03.439544	control	old_page
294476	697314	2017-01-15 01:20:28.957438	control	old_page
294477	715931	2017-01-16 12:40:24.467417	treatment	new_page

	converted	ab_page
0	0	0
1	0	0
2	0	1
3	0	1
4	1	0
...
294473	0	0
294474	0	0
294475	0	0
294476	0	0
294477	0	1

[290584 rows x 6 columns]

c. Use **statsmodels** to import your regression model. Instantiate the model, and fit the model using the two columns you created in part **b.** to predict whether or not an individual converts.

```
from scipy import stats
stats.chisqprob = lambda chisq, df2: stats.chi2.sf(chisq, df2)
```

```
df2['intercept'] = 1
```

```
lm = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
results = lm.fit()
```

```

Optimization terminated successfully.
      Current function value: 0.366118
      Iterations 6

```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
results.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```

=====
Dep. Variable:          converted    No. Observations:
290584
Model:                  Logit        Df Residuals:
290582
Method:                 MLE          Df Model:
1
Date:                  Fri, 09 Sep 2022    Pseudo R-squ.:
8.077e-06
Time:                  18:04:51    Log-Likelihood:
1.0639e+05
converged:              True          LL-Null:
1.0639e+05
Covariance Type:       nonrobust        LLR p-value:
0.1899
=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
intercept      -1.9888      0.008    -246.669      0.000      -2.005
-1.973
ab_page        -0.0150      0.011     -1.311      0.190      -0.037
0.007
=====
=====
"""

```

e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**? **Hint:** What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in the **Part II**?

p_valueu is 0.190.

Part II, the p-value was calculated where the null hypothesis was that the conversion rate for the new page more than the old page, and the alternative was the conversion rate for the old page less than or equal to the old page.

Part III, we used variables, and used a linear model to determine the p-value. The null hypothesis was that the difference between the pages = 0, and the alternative hypothesis was the difference between the pages $\neq 0$.

f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

well it is a good idea to add other factor to the model but it may be get more complexity and this is consider a disadvantage

g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives. You will need to read in the **countries.csv** dataset and merge together your datasets on the appropriate rows. [Here](#) are the docs for joining tables.

Does it appear that country had an impact on conversion? Don't forget to create dummy variables for these country columns - **Hint: You will need two columns for the three dummy variables**. Provide the statistical output as well as a written response to answer this question.

```
countries_df = pd.read_csv('./countries.csv')
df_new =
countries_df.set_index('user_id').join(df2.set_index('user_id'),
how='inner')
```

```
#show top rows from countries data
countries_df.head(10)
```

	user_id	country
0	834778	UK
1	928468	US
2	822059	UK
3	711597	UK
4	710616	UK
5	909908	UK
6	811617	US
7	938122	US
8	887018	US
9	820683	US

```
#show data in country column
countries_df['country'].unique()

array(['UK', 'US', 'CA'], dtype=object)
```

#show data after joining two files

`df_new.head(10)`

	country		timestamp	group	landing_page \
user_id					
834778	UK	2017-01-14	23:08:43.304998	control	old_page
928468	US	2017-01-23	14:44:16.387854	treatment	new_page
822059	UK	2017-01-16	14:04:14.719771	treatment	new_page
711597	UK	2017-01-22	03:14:24.763511	control	old_page
710616	UK	2017-01-16	13:14:44.000513	treatment	new_page
909908	UK	2017-01-06	20:44:26.334764	treatment	new_page
811617	US	2017-01-02	18:42:11.851370	treatment	new_page
938122	US	2017-01-10	09:32:08.222716	treatment	new_page
887018	US	2017-01-06	11:09:40.487196	treatment	new_page
820683	US	2017-01-14	11:52:06.521342	treatment	new_page

	converted	ab_page	intercept
user_id			
834778	0	0	1
928468	0	1	1
822059	1	1	1
711597	0	0	1
710616	0	1	1
909908	0	1	1
811617	1	1	1
938122	1	1	1
887018	0	1	1
820683	0	1	1

Create the necessary dummy variables

`df_new[['UK', 'US', 'CA']] = pd.get_dummies(df_new['country'])`

`df_new.head()`

	country		timestamp	group	landing_page \
user_id					
834778	UK	2017-01-14	23:08:43.304998	control	old_page
928468	US	2017-01-23	14:44:16.387854	treatment	new_page
822059	UK	2017-01-16	14:04:14.719771	treatment	new_page
711597	UK	2017-01-22	03:14:24.763511	control	old_page
710616	UK	2017-01-16	13:14:44.000513	treatment	new_page

	converted	ab_page	intercept	UK	US	CA
user_id						
834778	0	0	1	0	1	0
928468	0	1	1	0	0	1
822059	1	1	1	0	1	0
711597	0	0	1	0	1	0
710616	0	1	1	0	1	0

Fit Your Linear Model And Obtain the Results

`lm =`

```
sm.Logit(df_new['converted'],df_new[['intercept','ab_page','CA','US']]
)
results = lm.fit()
results.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.366113
      Iterations 6
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
=====
```

```
Dep. Variable:          converted    No. Observations:
290584
Model:                Logit        Df Residuals:
290580
Method:                MLE         Df Model:
3
Date:                Fri, 09 Sep 2022    Pseudo R-squ.:
2.323e-05
Time:                18:05:26    Log-Likelihood:    -
1.0639e+05
converged:            True        LL-Null:    -
1.0639e+05
Covariance Type:      nonrobust    LLR p-value:
0.1760
```

```
=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					

intercept	-2.0300	0.027	-76.249	0.000	-2.082
-1.978					
ab_page	-0.0149	0.011	-1.307	0.191	-0.037
0.007					
CA	0.0408	0.027	1.516	0.130	-0.012
0.093					
US	0.0506	0.028	1.784	0.074	-0.005
0.106					

```
=====
=====
"""
```

The p-value for ab_page is 0.191

The null hypothesis was that the difference in means would be 0, and the alternative was the difference would be greater or less than 0. The p-value is still large so that we fail to reject the null hypothesis.

then the difference in means would be 0

h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

#create interaction columns

```
df_new['us_ab_page'] = df_new.US *df_new.ab_page
df_new['ca_ab_page'] = df_new.CA *df_new.ab_page
```

```
lm =
sm.Logit(df_new['converted'],df_new[['intercept','ab_page','US','us_ab
_page','CA','ca_ab_page']])
results = lm.fit()
results.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.366109
      Iterations 6
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
=====
Dep. Variable:          converted    No. Observations:
290584
Model:                Logit        Df Residuals:
290578
Method:                MLE         Df Model:
5
Date:                  Fri, 09 Sep 2022    Pseudo R-squ.:
3.482e-05
Time:                  18:05:54           Log-Likelihood:    -
1.0639e+05
converged:              True            LL-Null:            -
1.0639e+05
Covariance Type:        nonrobust        LLR p-value:
0.1920
=====
```

```

=====
                                coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
intercept      -2.0040      0.036     -55.008      0.000      -2.075
-1.933
ab_page        -0.0674      0.052      -1.297      0.195      -0.169
0.034
US              0.0118      0.040       0.296      0.767      -0.066
0.090
us_ab_page      0.0783      0.057       1.378      0.168      -0.033
0.190
CA              0.0175      0.038       0.465      0.642      -0.056
0.091
ca_ab_page      0.0469      0.054       0.872      0.383      -0.059
0.152
=====
=====
" " "

```

Summary and conclusion on regression

p_value for both interactions > 0.05.

there is no difference between influence of landing_page in US and other countries.

Conclusions

there is not enough evidence that the new_page increases the conversion rate as compared to the old_page.

no evidence that the countries [US, UK ,CA] influence the conversion rate.